

1 # Heart Disease Prediction using Machine Learning

1 importing the Libraries

In [9]:

```
1 import numpy as np
2 import pandas as pd
3 from sklearn.model_selection import train_test_split
4 from sklearn.linear_model import LogisticRegression
5 from sklearn.metrics import accuracy_score
```

1 Data collection and processing

In [10]:

```
1 #Loading the dataset to a pandas Dataframe
2 heart_data = pd.read_csv("C:\\\\Users\\\\trylogic\\\\Downloads\\\\heart.csv")
```

In [11]:

```
1
2 # print the first five rows
3 heart_data.head()
```

Out[11]:

	age	sex	cp	trestbps	chol	fb	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3	0
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3	0
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3	0
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3	0
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2	0

◀ ▶

In [12]:

```
1 # print Last five rows
2 heart_data.tail()
```

Out[12]:

	age	sex	cp	trestbps	chol	fb	restecg	thalach	exang	oldpeak	slope	ca	thal	target
1020	59	1	1	140	221	0	1	164	1	0.0	2	0	2	
1021	60	1	0	125	258	0	0	141	1	2.8	1	1	3	
1022	47	1	0	110	275	0	0	118	1	1.0	1	1	2	
1023	50	0	0	110	254	0	0	159	0	0.0	2	0	2	
1024	54	1	0	120	188	0	1	113	0	1.4	1	1	3	

◀ ▶

In [13]:

```
1 # number of rows and columns
2 heart_data.shape
```

Out[13]:

(1025, 14)

In [14]:

```
1 heart_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   age         1025 non-null   int64  
 1   sex          1025 non-null   int64  
 2   cp           1025 non-null   int64  
 3   trestbps    1025 non-null   int64  
 4   chol         1025 non-null   int64  
 5   fbs          1025 non-null   int64  
 6   restecg     1025 non-null   int64  
 7   thalach     1025 non-null   int64  
 8   exang        1025 non-null   int64  
 9   oldpeak      1025 non-null   float64 
 10  slope        1025 non-null   int64  
 11  ca           1025 non-null   int64  
 12  thal         1025 non-null   int64  
 13  target       1025 non-null   int64  
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

In [15]:

```
1 # checking for missing values
2 heart_data.isnull().sum()
```

Out[15]:

```
age          0
sex          0
cp           0
trestbps    0
chol         0
fbs          0
restecg     0
thalach     0
exang        0
oldpeak      0
slope        0
ca           0
thal         0
target       0
dtype: int64
```

In [18]:

```
1 heart_data.describe() #describe
```

Out[18]:

	age	sex	cp	trestbps	chol	fbs	reste
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000
mean	54.434146	0.695610	0.942439	131.611707	246.000000	0.149268	0.5297
std	9.072290	0.460373	1.029641	17.516718	51.59251	0.356527	0.5278
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.0000
25%	48.000000	0.000000	0.000000	120.000000	211.000000	0.000000	0.0000
50%	56.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.0000
75%	61.000000	1.000000	2.000000	140.000000	275.000000	0.000000	1.0000
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.0000

In [19]:

```
1 # checking the distribution of target variables
2 heart_data['target'].value_counts()
```

Out[19]:

```
1      526
0      499
Name: target, dtype: int64
```

```
1 1 ==> defective heart
2 0 ==> is healthy heart
```

```
1 Splitting Features and Target
```

In [20]:

```
1 X= heart_data.drop(columns='target',axis=1)
2 Y=heart_data['target']
```

In [21]:

```
1 print(X)
2
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	52	1	0	125	212	0	1	168	0	1.0	
1	53	1	0	140	203	1	0	155	1	3.1	
2	70	1	0	145	174	0	1	125	1	2.6	
3	61	1	0	148	203	0	1	161	0	0.0	
4	62	0	0	138	294	1	1	106	0	1.9	
...
1020	59	1	1	140	221	0	1	164	1	0.0	
1021	60	1	0	125	258	0	0	141	1	2.8	
1022	47	1	0	110	275	0	0	118	1	1.0	
1023	50	0	0	110	254	0	0	159	0	0.0	
1024	54	1	0	120	188	0	1	113	0	1.4	
	slope	ca	thal								
0	2	2	3								
1	0	0	3								
2	0	0	3								
3	2	1	3								
4	1	3	2								
...								
1020	2	0	2								
1021	1	1	3								
1022	1	1	2								
1023	2	0	2								
1024	1	1	3								

[1025 rows x 13 columns]

In [22]:

```
1 print(Y)
```

0	0
1	0
2	0
3	0
4	0
..	..
1020	1
1021	0
1022	0
1023	1
1024	0

Name: target, Length: 1025, dtype: int64

```
1 Splitting the data into training data and test data
```

In [23]:

```
1 X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.2,stratify=Y,random_st
```

In [24]:

```
1 print(X.shape,X_train.shape,X_test.shape)
```

(1025, 13) (820, 13) (205, 13)

In [25]:

```
1 print(X_train)
2 print(Y_train)
```

	age	sex	cp	trestbps	chol	fb	restecg	thalach	exang	oldpeak	\
148	52	1	3	152	298	1	1	178	0	1.2	
493	55	1	0	132	353	0	1	132	1	1.2	
991	60	1	0	117	230	1	1	160	1	1.4	
52	38	1	2	138	175	0	1	173	0	0.0	
210	42	1	2	120	240	1	1	194	0	0.8	
..
992	50	0	0	110	254	0	0	159	0	0.0	
601	46	1	0	140	311	0	1	120	1	1.8	
356	59	1	0	164	176	1	0	90	0	1.0	
747	60	1	0	117	230	1	1	160	1	1.4	
572	34	1	3	118	182	0	0	174	0	0.0	
	slope	ca	thal								
148	1	0	3								
493	1	1	3								
991	2	2	3								
52	2	4	2								
210	0	0	3								
..								
992	2	0	2								
601	1	2	3								
356	1	2	1								
747	2	2	3								
572	2	0	2								

[820 rows x 13 columns]

	target
148	1
493	0
991	0
52	1
210	1
..	
992	1
601	0
356	0
747	0
572	1

Name: target, Length: 820, dtype: int64

```
1 Model Training
```

```
1 logistic regression
```

In [26]:

```
1 model = LogisticRegression()
```

In [27]:

```
1 # train the Logistic regression model
2 model.fit(X_train,Y_train)
```

C:\Users\trylogic\AppData\Local\Programs\Python\Python39\lib\site-packages\s
klearn\linear_model_logistic.py:814: ConvergenceWarning: lbfgs failed to co
nverge (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scik>
it-learn.org/stable/modules/preprocessing.html)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regre
ssion (https://scikit-learn.org/stable/modules/linear_model.html#logistic-re
gression)
n_iter_i = _check_optimize_result(

Out[27]:

LogisticRegression()

```
1 model evaluation and accuracy score
```

In [28]:

```
1 #accuracy on training data
2 X_train_prediction = model.predict(X_train)
3 training_data_accuracy = accuracy_score(X_train_prediction,Y_train)
```

In [29]:

```
1 print('Accuracy on training data : ', training_data_accuracy)
```

Accuracy on training data : 0.85

In [30]:

```
1 #accuracy on test data
2 X_test_prediction = model.predict(X_test)
3 test_data_accuracy = accuracy_score(X_test_prediction, Y_test)
```

In [31]:

```
1 print('Accuracy on test data : ', test_data_accuracy)
```

Accuracy on test data : 0.8390243902439024

```
1 building a predictive system
```

In [32]:

```
1 input_data = (62,0,0,138,294,1,1,106,0,1.9,1,3,2)
2
3 # changing the input_data to a numpy array
4 input_data_as_numpy_array = np.asarray(input_data)
5
6 # reshape the np array as we are predicting for one instance
7 input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)
8
9 prediction = model.predict(input_data_reshaped)
10 print(prediction)
11
12 if (prediction[0]== 0):
13     print('The person does not have a heart disease')
14 else:
15     print('The person has heart disease')
```

[0]

The person does not have a heart disease

```
C:\Users\trylogic\AppData\Local\Programs\Python\Python39\lib\site-packages\s
klearn\base.py:445: UserWarning: X does not have valid feature names, but Lo
gisticRegression was fitted with feature names
    warnings.warn(
```

In []:

1

In []:

1